

An Iterated Greedy heuristic for the sequence dependent setup times flowshop problem with makespan and weighted tardiness objectives

Rubén Ruiz ^{a,*}, Thomas Stützle ^b

^a *Universidad Politécnica de Valencia, Departamento de Estadística e Investigación, Operativa Aplicadas y Calidad, Grupo de Investigación Operativa, Camino de Vera S/N, 46021 Valencia, Spain*

^b *Université Libre de Bruxelles, IRIDIA, CP 194/6, Avenue Franklin Roosevelt 50, 1050 Bruxelles, Belgium*

Received 19 January 2006; accepted 18 July 2006

Available online 3 November 2006

Abstract

Iterated Greedy (IG) algorithms are based on a very simple principle, are easy to implement and can show excellent performance. In this paper, we propose two new IG algorithms for a complex flowshop problem that results from the consideration of sequence dependent setup times on machines, a characteristic that is often found in industrial settings. The first IG algorithm is a straightforward adaption of the IG principle, while the second incorporates a simple descent local search. Furthermore, we consider two different optimization objectives, the minimization of the maximum completion time or makespan and the minimization of the total weighted tardiness. Extensive experiments and statistical analyses demonstrate that, despite their simplicity, the IG algorithms are new state-of-the-art methods for both objectives.

© 2006 Elsevier B.V. All rights reserved.

Keywords: Iterated Greedy; Flowshop scheduling; Sequence dependent setup times; Makespan; Weighted tardiness; Stochastic local search; Metaheuristics

1. Introduction

In the flowshop problem we are given a set $N = \{1, \dots, n\}$ of n independent jobs that have to be processed on a set $M = \{M_1, \dots, M_m\}$ of m machines. Without loss of generality, the machines are arranged in order so that all jobs are processed

sequentially, i.e., first on machine 1, then on machine 2 and so on until the last machine m , which gives a total of $n \times m$ operations. Each operation requires a known and deterministic fixed amount of time that is commonly denoted as p_{ij} , $i \in M$, $j \in N$.

In the flowshop literature, one can find an overwhelming number of papers for the regular flowshop problem with the objective of minimizing the maximum completion time across all jobs (also called makespan and denoted by C_{\max}). However, the Sequence Dependent Setup Time Flowshop

* Corresponding author. Tel.: +34 96 387 7007; fax: +34 96 387 7499.

E-mail addresses: r Ruiz@eio.upv.es (R. Ruiz), stuetzle@ulb.ac.be (T. Stützle).

Problem (SDST flowshop problem in short) has attracted much less attention. Specific to the SDST flowshop are the setup times. Setup times involve operations that have to be performed on machines and that are not part of the processing times. This includes cleaning, fixing or releasing parts to machines, adjustments to machines, etc. The most complex situation arises when the setup times are separable from the processing time (i.e., the setup operations can be done before the product arrives to the machine) and when they are sequence dependent (i.e., the amount of setup time depends on the machine, on the job that the machine was processing and on the job that comes next). An example for the SDST flowshop comes from the paint industry; after producing a black paint, substantial cleaning must be performed if one intends to produce white paint, while less cleaning is necessary if a batch of dark grey paint is to be produced. In this paper, we consider this last type of setup times and denote by S_{ijk} the known, deterministic and non-negative setup time on machine i when producing job k , $k \in N$, after having produced job j .

The objective in flowshop scheduling problems is to find a sequence for processing the jobs on the machines so that a given criterion is optimized. This yields a total of $n!$ possible orderings of the operations on each machine and a total of $(n!)^m$ possible processing sequences. In flowshop scheduling research usually only so-called permutation sequences are considered, where the processing order of operations is the same for all machines. Here, we also adopt this restriction. In this paper, we consider two independent objectives, namely the minimization of the makespan and the minimization of the total weighted tardiness. According to Pinedo (2002), the two problems are denoted as $F/S_{ijk}, pmu/C_{\max}$ and $F/S_{ijk}, pmu/\sum_{j=1}^n w_j T_j$, respectively; in the following, we refer to these two problems as SDST-FSP- C_{\max} and SDST-FSP-WT, respectively. Makespan is one of the most commonly tackled objectives. In the pure flowshop context, minimizing the maximum completion time is equal to minimizing the idle time of machines. This objective is also closely related to the maximization of the throughput, which is very important in production environments where machines are expensive to purchase and operate. However, the best sequence with respect to makespan minimization might have a large number of jobs being completed after their due dates. Hence, we also consider the minimization of the weighted tardiness. With

this second, independent objective, every job has a weight or priority that indicates its relative importance. This priority is multiplied by the amount of time the job is completed after its due date. Weighted tardiness minimization is of uttermost importance in make-to-order or just-in-time environments, where due dates are typically set by clients.

Regarding the computational complexity, the SDST flowshop with the C_{\max} objective has been shown to be \mathcal{NP} -hard by Gupta (1986) even when $m = 1$ and also when $m = 2$ and setups are present only on the first or second machine (see Gupta and Darrow, 1986). For $m = 1$, the SDST flowshop is known to be a special case of the Traveling Salesman Problem (TSP) that is also well known to be \mathcal{NP} -hard. Based on the complexity hierarchies for scheduling problems (Pinedo, 2002), we can deduce that also the SDST flowshop with total weighted tardiness objective is \mathcal{NP} -hard.

Recently, a rather new heuristic method, called Iterated Greedy (IG), has shown state-of-the-art performance for the permutation flowshop scheduling problem where the goal is to minimize the makespan (FSP- C_{\max}) in Ruiz and Stützle (2007). This is rather noteworthy since a large number of often rather complex and very fine-tuned algorithms have been proposed for the FSP- C_{\max} (Ruiz and Maroto, 2005) and the IG method is remarkably simple. IG for the FSP- C_{\max} works by iteratively applying two phases; in a *destruction* phase, some jobs are eliminated from the current solution, and in the *construction* phase the eliminated jobs are reinserted into the sequence using a greedy construction heuristic. To these two phases, an additional local search phase can easily be added.

The main goal of this article is an experimental study as to whether the high performance of IG for the FSP- C_{\max} is transferable also to other, more complex flowshop environments, in particular SDST flowshop problems. In fact, the results obtained with the proposed IG algorithms show that this is actually the case and that the IG algorithm with an additional local search phase is a new state-of-the-art algorithm for the SDST-FSP- C_{\max} and the SDST-FSP-WT.

The remainder of the paper is organized as follows: Section 2 reviews the literature on the two problems considered here. In Section 3, we introduce the two IG algorithms proposed in this paper. A full experimental evaluation and comparison of the proposed methods is shown, along with statistical

analyses, in Section 4. Finally, Section 5 closes this study with some conclusions and future research directions.

2. Literature review

Compared to the regular flowshop, on which hundreds of papers have been published, the literature on the SDST counterpart is scarce. In a recent article, Ruiz et al. (2005) carried out an extensive literature survey about this problem and here we summarize the most important aspects.

Exact techniques for the SDST flowshop have shown rather limited results. The works of Corwin and Esogbue (1974) on dynamic programming or Srikar and Ghosh (1986), Stafford and Tseng (1990), Tseng and Stafford (2001) and Stafford and Tseng (2002) on integer programming or the papers of Ríos-Mercado and Bard (1998a), Ríos-Mercado and Bard (1999a) and Ríos-Mercado and Bard (2003) on branch and bound and related exact methods are able to optimally solve problem instances with only about 10 jobs and few machines in general and in most cases only the C_{\max} objective is considered. Some heuristics and applications of stochastic local search algorithms to the SDST-FSP- C_{\max} have also been proposed. Simons (1992) proposed two general heuristics for the SDST flowshop, named TOTAL and SETUP. Das et al. (1995) proposed a heuristic method based on a savings index. Ríos-Mercado and Bard (1998b) proposed a modification of the well-known NEH heuristic for the regular flowshop from Nawaz et al. (1983) that considered setup times and that they called NEH_RMB; in the same article they proposed a GRASP algorithm. In a later work, the same authors proposed a modification of the heuristics of Simons (1992) resulting in a new heuristic called HYBRID (Ríos-Mercado and Bard, 1999b). Recently, Ruiz et al. (2005) proposed a genetic and a memetic algorithm for the SDST-FSP- C_{\max} . They carried out an extensive experimental study and compared these two algorithms to all aforementioned algorithms as well as to many adaptations of other methods that were proposed for the FSP- C_{\max} . The result was that both algorithms, especially the memetic algorithm, are clearly superior to all other alternatives.

On the SDST-FSP-WT, little has been published. In two similar articles (Parthasarathy and Rajendran, 1997a,b), a Simulated Annealing heuristic was proposed for the SDST flowshop problem with

the objectives of minimizing the maximum weighted tardiness and the total weighted tardiness. In a more recent work, Rajendran and Ziegler (2003) introduced a new heuristic paired with a local search improvement scheme for a combined objective of total weighted flow-time and tardiness. Another similar work is Rajendran and Ziegler (1997) where only weighted flow-time is considered.

Allahverdi et al. (in press) have put together a much more updated and comprehensive review of scheduling research with setup times in which no other relevant papers related to the SDST flowshop can be found.

3. IG for the SDST Flowshop

Iterated Greedy algorithms start from some initial solution and then iterate through a main loop in which first a partial candidate solution s_p is obtained by removing a number of solution components from a complete candidate solution s and next a complete solution s' is re-constructed starting with s_p . Before continuing with the next loop, an acceptance criterion decides whether s' becomes the new incumbent solution. This process is repeated until some stopping criterion like a maximum number of iterations or a computation time limit is met. It is also straightforward to add an optional local search phase for improving the re-constructed solution s' before applying the acceptance test. An outline of a generic IG algorithm is given in Fig. 1.

IG is related to other stochastic local search (SLS) methods (see Hoos and Stützle, 2004), and especially to Iterated Local Search (ILS) (see Lourenço et al., 2003).¹ The applications of the IG method are, for the moment rather few. Examples of applications of IG include the researches of Jacobs and Brusco (1995) and Marchiori et al. (2000) on the Set Covering Problem, the work of Cesta et al. (2000) and Laurent and Van Hentenryck (2004) on some scheduling problems, and the recent application of two simple IG algorithms to the FSP- C_{\max} (Ruiz and Stützle, 2007). In this latter work,

¹ One main difference between IG and ILS is that while the central idea of IG is to iterate over constructive algorithms, the central idea of ILS resides in iterating over local searches in a specific way. Nevertheless, especially when adding a local search phase to IG algorithms, as explained later on, the similarities between IG and ILS become even more pronounced. Note that occasionally IG has also been referred to as Ruin-and-Recreate (Schrimpf et al., 2000) or iterative flattening (Cesta et al., 2000) and (Laurent and Van Hentenryck, 2004).

```

procedure Iterated_Greedy
   $s_0 := \text{GenerateInitialSolution};$ 
   $s := \text{LocalSearch}(s_0);$            %optional
  repeat
     $s_p := \text{Destruction}(s)$ 
     $s' := \text{Construction}(s_p)$ 
     $s' := \text{LocalSearch}(s')$        %optional
     $s := \text{AcceptanceCriterion}(s, s')$ 
  until termination condition met
end

```

Fig. 1. General outline of an Iterated Greedy algorithm.

the simple IG algorithm yielded results far better than several other, more complex approaches.

Before entering into a detailed description of the IG algorithms for the two SDST flowshop problems tackled here, we give additional notation and describe how to calculate the C_{\max} and $\sum_{j=1}^n w_j T_j$ objectives of a given sequence. For every job j , $j \in N$ we have a due date denoted as d_j and a weight w_j . We denote with $C_{i,j}$ the completion time of job j on machine i . The completion time of job j at the last stage m is $C_{m,j}$ or C_j in short. The *lateness* of job j is defined as $L_j = C_j - d_j$. Note that L_j can be positive or negative and, hence, the *tardiness* of a job j is defined as $T_j = \max\{L_j, 0\}$.

Let π be a permutation or sequence of the n jobs. The job in position j of the sequence is denoted as $\pi(j)$. The completion times for the jobs on the machines are calculated with the formula:

$$C_{i,\pi(j)} = \max\{C_{i,\pi(j-1)} + S_{i,\pi(j-1),\pi(j)}; C_{i-1,\pi(j)}\} + p_{i,\pi(j)}, \quad (1)$$

where $C_{0,\pi(j)} = 0$, $S_{i,0,\pi(j)} = 0$ and $C_{i,0} = 0$, for all $i \in M$, $j \in N$. With this we have that $C_{\max} = C_{\pi(n)}$ and with the completion times on the last machine for all jobs we can calculate the total weighted tardiness objective as $\sum_{j=1}^n w_j T_j$.

3.1. Algorithm initialization

The initial solution for IG is ideally generated by a high performance construction heuristic. For the SDST-FSP- C_{\max} , we used the NEH_RMB heuristic proposed by Ríos-Mercado and Bard (1998b); it is an extension of the well-known NEH heuristic from Nawaz et al. (1983). NEH_RMB has a computational complexity of $\mathcal{O}(n^2m)$, since for the SDST-FSP- C_{\max} the same speed-ups as proposed by Taillard (1990) for the NEH heuristic can be applied.

Recall that NEH is an insertion heuristic, where at each step the next unscheduled job is tentatively inserted in each possible position of some partial solution. The job is then finally inserted into the position where the objective function takes the lowest value. For executing such an insertion heuristic, the jobs need to be ordered in some way. For more details how this is done in NEH_RMB, the reader is referred to Ríos-Mercado and Bard (1998b).

For the SDST-FSP-WT, the situation is rather different. As we will see in later sections, there are no known high performing construction heuristics available. Hence, we extended the original NEH, which has been widely used for many other scheduling problems, to this second problem, resulting in a heuristic we call NEH_EWDD. However, the accelerations proposed by Taillard (1990) are not transferable for objectives other than C_{\max} and, hence, the complexity of NEH_EWDD is $\mathcal{O}(n^3m)$. In NEH_EWDD we consider the existence of weights and due dates for defining an initial order in which the jobs are considered for insertion. The initial order in NEH_EWDD is based on the Earliest Weighted Due Date dispatching rule (EWDD) that arranges jobs in ascending order of their weighted due dates, i.e., d_j/w_j . We carried out some experiments with other initialization procedures such as the HA algorithm of Rajendran and Ziegler (1997, 2003), which will be evaluated later, and found no improvements. Therefore, we choose of NEH_EWDD for initialization for this second objective.

3.2. Destruction, construction and acceptance criterion

As said above, once the initial sequence has been obtained, the destruction step is applied. Let the initial sequence be π . In this step, a given number d of jobs, chosen at random and without repetition, are removed from the sequence. This creates two subsequences, the first one of size d is called π_R and contains the removed jobs in the order in which they were removed. The second, of size $n - d$, is the original one without the removed jobs; this sequence we call π_D . If $d < n$ and constant, the complexity of this phase is $\mathcal{O}(1)$, if a permutation is represented via a linked list.

After the two subsequences are created, the construction step is applied. Here, all jobs in π_R are reinserted in π_D by applying the NEH heuristic, i.e., the first job of π_R ($\pi_{R(1)}$) is inserted in all possible $n - d + 1$ positions of π_D generating $n - d + 1$

partial sequences that include job $\pi_{R(1)}$. The sequence with the best objective value is kept for the following iteration. The process ends when π_R is empty and therefore π_D is again of size n . It has to be noted that for the SDST-FSP- C_{\max} we use in this construction phase the speedier NEH_RMB heuristic, which results in a complexity for this construction phase of $\mathcal{O}(dnm)$. The computational complexity for weighted tardiness criterion is $\mathcal{O}(dn^2m)$.

Another step in the Iterated Greedy algorithm is to decide whether the reconstructed sequence is accepted or not as the incumbent solution for the next iteration. A pure descent criterion would be to accept solutions with better objective function values. However, this acceptance criterion is prone to stagnation. As an alternative, we consider an acceptance criterion that is frequently used in simulated annealing algorithms but with a constant temperature value. This acceptance criterion has been used in other papers like in Osman and Potts (1989), Stützle (1998), or Ruiz and Stützle (2007) and there the value of Temperature depended on the processing times, the number of jobs and machines, and another adjustable parameter T :

$$\text{Temperature} = T \frac{\sum_{i=1}^m \sum_{j=1}^n p_{ij}}{n \times m \times 10}. \quad (2)$$

We denote by $F(\text{incumbent})$ and $F(\pi)$ the respective objective function values of the incumbent sequence and the reconstructed sequence. If $F(\pi) \geq F(\text{incumbent})$, π is accepted as the new incumbent solution if

$$\text{random} \leq \exp\{F(\text{incumbent}) - F(\pi) / \text{Temperature}\},$$

where random is a random number uniformly distributed in $[0, 1]$.

```

procedure LocalSearch_Insertion ( $\pi$ )
    improve := true;
    while (improve = true) do
        improve := false;
        for  $i := 1$  to  $n$  do
            remove a job  $k$  at random from  $\pi$  (without repetition)
             $\pi' :=$  best permutation obtained by inserting  $k$  in all possible positions of  $\pi$ ;
            if  $F(\pi') < F(\pi)$  then
                 $\pi := \pi'$ ;
                improve := true;
            endif
        endfor
    endwhile
    return  $\pi$ 
end

```

Fig. 2. Descent local search procedure based on the insertion neighborhood.

3.3. Hybridization with local search

To further improve the performance of the basic IG algorithm, Ruiz and Stützle (2007) added a simple descent local search phase after the NEH initialization and after each sequence re-construction. In fact, such a hybridization of construction heuristics with local search is common practice in stochastic local search. We apply a local search based on the insertion neighborhood. This neighborhood has been regarded as superior to the swap or exchange neighborhoods in flowshop scheduling by many authors (see Osman and Potts, 1989; Taillard, 1990; Stützle, 1998 or Ruiz et al., 2005 among others). The insertion neighborhood of a sequence comprises all those sequences that can be obtained by removing some job and inserting it in another position. We implemented a local search algorithm that searches this neighborhood in a particular way: In each local search step, a job is removed from the sequence (at random and without repetition) and then inserted in all possible n positions. The current sequence π is replaced by the best sequence among the n possible ones, only if an improvement of π can be obtained. The procedure ends when no further improvements are found, i.e., when the sequence is a local optimum with respect to the insertion neighborhood. Since this local search is applied until a local optimum is found, we can only derive the computational complexity of one single pass. This complexity is $\mathcal{O}(n^2m)$ and $\mathcal{O}(n^3m)$ for C_{\max} and total weighted tardiness objectives, respectively. The pseudocode of this procedure is shown in Fig. 2.

With this optional local search step, the proposed IG algorithm is depicted in Fig. 3.


```

procedure IteratedGreedy_for_SDST_flowshop
   $\pi := \text{NEH\_RMB}$  or  $\text{NEH\_EWDD}$ ;           % Initialization
   $\pi := \text{LocalSearch\_Insertion}(\pi)$ ;
   $\pi_b := \pi$ ;
  while (termination criterion not satisfied) do
     $\pi' := \pi$ ;                               % Destruction phase
    for  $i := 1$  to  $d$  do
       $\pi' :=$  remove one job at random from  $\pi'$  and insert it in  $\pi'_R$ ;
    endfor
    for  $i := 1$  to  $d$  do                       % Construction phase
       $\pi' :=$  best permutation obtained by inserting job  $\pi_{R(i)}$  in all possible positions of  $\pi'$ ;
    endfor
     $\pi'' := \text{LocalSearch\_Insertion}(\pi')$ ; % Local Search
    if  $F(\pi'') < F(\pi)$  then                 % Acceptance Criterion
       $\pi := \pi''$ ;
      if  $F(\pi) < F(\pi_b)$  then             % check if new best permutation
         $\pi_b := \pi$ ;
      endif
    elseif ( $\text{random} \leq \exp\{-(F(\pi'') - F(\pi))/\text{Temperature}\}$ ) then
       $\pi := \pi''$ ;
    endif
  endwhile
  return  $\pi_b$ 
end

```

Fig. 3. Iterated Greedy algorithm with the optional local search phase. See the text for more details.

4. Experimental evaluation

We refer to the basic IG algorithm as IG_RS and as IG_RS_{LS} to the one including the local search phase. For both we set $T = 0.5$ and $d = 4$. These parameter settings follow the extensive experimental study carried out in Ruiz and Stützle (2007). We tested in some preliminary experiments other settings for the two parameters without finding any strong improvements. Therefore, we keep the original values as stated. Next, we present the results of both algorithms and compare them to other existing methods for each of the two objectives separately.

4.1. Experimental results for the SDST-FSP- C_{\max}

Ruiz et al. (2005) generated four instance sets for testing algorithms for the SDST-FSP- C_{\max} . Each set is based on the original 120 instances of Taillard (1993), which are organized in 12 groups with 10 instances each. The groups contain different combinations of the number of jobs n and the number of machines m . The $n \times m$ combinations are: $\{20, 50, 100\} \times \{5, 10, 20\}$, $200 \times \{10, 20\}$ and 500×20 . The processing times (p_{ij}) in Taillard's instances are generated from a uniform distribution in the range $[1, 99]$. To generate the four instance sets for the SDST-FSP- C_{\max} , the ratios of the setup

times to the processing times were changed such that the sequence dependent setup times are at most 10%, 50%, 100% or 125%, respectively, of the maximum possible processing times of Taillard's original instances. This results in four sets of instances called SDST10, SDST50, SDST100, and SDST125 that have the setup times uniformly distributed in the range $[1, 9]$, $[1, 49]$, $[1, 99]$ and $[1, 124]$, respectively. Thus, we have a total of 480 different instances, which can be downloaded from <http://www.upv.es/gio/rruiz>.

We compare a total of six algorithms. The first two are the proposed IG_RS and IG_RS_{LS} algorithms. We also test the genetic and memetic algorithms of Ruiz et al. (2005). We refer to the genetic algorithm as GA. The memetic algorithm (MA) in the original work used a descent local search that essentially consisted in one execution of the while loop of the descent local search given in Fig. 2 and, hence, it does not necessarily return a locally optimal solution. This choice was done to reduce the computation time of the local search and to allow for more iterations of the MA in the same computation time. In this article, we propose a modified version that uses the same local search as IG_RS_{LS} in the local search phase of the memetic algorithm. We refer to the modified memetic algorithm as MA_{LS}. The last algorithm that we test is

the PACO ant colony optimization algorithm from Rajendran and Ziegler (2004). PACO was adapted to the SDST-FSP- C_{\max} and the initialization was changed from the original NEH to NEH_RMB. It has to be noted that we could have tested also other SLS algorithms that were proposed for the SDST flowshop (see Section 2 for details). However, from the extensive comparison in Ruiz et al. (2005), which tested a total of 14 methods, the memetic algorithm (denoted here as MA) resulted to be clearly superior to all others. Consequently, here we only compare to that algorithm and to the others aforementioned.

For the experiments we use a PC/AT computer with an Athlon XP 1600+ processor (1400 MHz) and 512 MBytes of main memory. Every algorithm is run 10 independent times with a stopping criterion based on an elapsed CPU time given by $(n \times m/2) \times t$ milliseconds. This allows for more time as the number of jobs n and the number of machines m grows. We report experimental results for three settings of t to 30, 60 and 90, respectively; this allows us to examine the relative ranking of the algorithms in dependence of the available computation time. The longest computation times in these experiments are for the largest instances (500×20) and for $t = 90$, where each trial of an algorithm takes $(500 \times 20/2) \times 90 = 450,000$ milliseconds or 7.5 minutes. The response variable is the relative percentage deviation over the best known solution and it is calculated as

Relative percentage deviation (RPD)

$$= \frac{\text{Some}_{\text{sol}} - \text{Best}_{\text{sol}}}{\text{Best}_{\text{sol}}} \times 100, \quad (3)$$

where Some_{sol} is the solution returned for a given instance and Best_{sol} is the best known solution for the same instance. The best known solutions for the four instance sets can be downloaded from <http://www.upv.es/gio/rruiz>. These best known solutions are the best ones identified in previous works (see Ruiz et al., 2005) and in this article. Another option would have been to give average deviations from lower bounds. However, we refrained from doing so, since the available lower bounds for the SDST-FSP- C_{\max} are rather weak. For example, on instances of size 20×6 , which are of similar size as the smallest ones treated in this article, the branch and bound algorithm of Ríos-Mercado and Bard (1999a) left optimality gaps of ca. 8% after 30 minutes of computation time. Later work like that presented in Ríos-Mercado and Bard

(2003), apparently does not deliver much tighter lower bounds.

The results for instance sets SDST10 and SDST50 are shown in Table 1 and for instance sets SDST100 and SDST125 in Table 2. For each combination of n and m , these results are averaged across the 10 instances and across all repetitions; hence, each cell of the tables gives the average percentage deviation for each value of t (i.e., three average results, one for $t = 30, 60$ and 90) across $10 \times 10 = 100$ values. We now first discuss the results based on the average results and we later comment, with thorough statistical analyses, on the statistical significance of the differences. (In a nutshell, the later statistical analysis shows that almost all differences observed under same settings of the parameter t are significant.)

From the cross averages in the tables we can see that each of the algorithms can profit from additional computation time, since with increasing value of parameter t , also the average relative percentage deviation is reduced. Additionally, the differences in the solution quality reached by the algorithms can be rather high. It is convenient to consider also the ranking of the algorithms on the instance sets. If we consider the cross averages in the tables (last line for each instance set), we can see that the overall ranking (ranking from best to highest cross averages) of the algorithms remains very similar across all four sets of instances and also across the different settings of t . In fact, the overall ranking would be the same for the sets SDST50, SDST100, and SDST125 and it is IG_RS_{LS}, MA_{LS}, MA, IG_RS, PACO, and GA from best performing to worst performing algorithm. Only on instance set SDST10, PACO would rank as the third best algorithm overall with the others maintaining their same relative ordering as on the other three instance sets.

These results suggest that IG_RS_{LS} is clearly the best performing algorithm independent of the computation time limit chosen and of the magnitude of the sequence dependent setup times. This is remarkable, given the simplicity of the algorithm, when compared to, for example, the much more complex memetic algorithm MA_{LS} that requires to maintain a population of solutions and also more complex operators like recombination among solutions. IG_RS_{LS} and MA_{LS} use the very same descent local search algorithm. Hence, the guiding mechanism for the search used in the IG algorithm appears here to be the main reason for the difference in performance.

Table 1

Average relative percentage deviation from the best known solutions for the algorithms for the SDST-FSP- C_{\max} on instance sets SDST10 and SDST50 with the termination criteria set at $(n \times m/2) \times t$ milliseconds maximum CPU time

Instance	GA	MA	MA _{LS}	PACO	IG_RS	IG_RS _{LS}
<i>SDST10 instances</i>						
20 × 5	0.43/0.46/0.41	0.49/0.90/0.70	0.12/0.10/0.08	0.18/0.21/0.18	0.21/0.19/0.14	0.08/0.05/0.04
20 × 10	0.59/0.57/0.56	0.55/0.28/0.36	0.13/0.13/0.13	0.33/0.26/0.22	0.28/0.22/0.24	0.08/0.05/0.04
20 × 20	0.44/0.37/0.39	0.59/0.52/0.56	0.14/0.09/0.10	0.20/0.16/0.12	0.30/0.22/0.19	0.07/0.05/0.04
50 × 5	1.04/0.93/0.92	0.77/0.57/0.77	0.43/0.31/0.30	0.53/0.44/0.42	1.00/0.88/0.84	0.37/0.32/0.27
50 × 10	2.10/2.07/2.01	1.21/1.38/1.26	1.12/0.83/0.81	1.23/1.02/1.06	1.58/1.58/1.43	0.76/0.60/0.53
50 × 20	2.23/2.18/2.10	1.38/1.21/1.28	1.16/0.96/0.82	1.27/1.06/1.01	1.85/1.70/1.54	0.91/0.64/0.60
100 × 5	1.28/1.10/1.03	0.76/0.70/0.63	0.54/0.40/0.31	0.87/0.80/0.76	1.44/1.36/1.34	0.43/0.38/0.33
100 × 10	1.48/1.39/1.33	0.91/0.81/0.90	0.78/0.60/0.48	0.99/0.84/0.77	1.49/1.37/1.32	0.61/0.44/0.38
100 × 20	2.07/1.93/1.83	1.49/1.11/1.06	1.27/0.97/0.82	1.49/1.25/1.12	1.75/1.48/1.47	0.88/0.71/0.54
200 × 10	1.63/1.42/1.32	0.81/0.73/0.65	0.79/0.61/0.48	1.04/0.94/0.85	1.50/1.39/1.33	0.58/0.43/0.32
200 × 20	2.00/1.79/1.71	1.14/0.93/0.87	1.11/0.87/0.76	1.31/1.10/0.95	1.45/1.25/1.12	0.79/0.53/0.38
500 × 20	1.38/1.31/1.27	0.74/0.54/0.48	0.69/0.54/0.43	0.82/0.69/0.61	1.01/0.88/0.82	0.46/0.31/0.21
Average	1.39/1.29/1.24	0.90/0.81/0.79	0.69/0.53/0.46	0.86/0.73/0.67	1.16/1.04/0.98	0.50/0.38/0.31
<i>SDST50 instances</i>						
20 × 5	1.34/1.30/1.15	0.44/1.21/1.50	0.37/0.35/0.30	0.58/0.53/0.51	0.83/0.69/0.58	0.26/0.18/0.10
20 × 10	1.21/1.16/1.17	0.92/0.87/0.77	0.41/0.31/0.32	0.49/0.43/0.44	0.66/0.62/0.58	0.28/0.20/0.19
20 × 20	0.57/0.57/0.49	0.87/0.23/0.78	0.20/0.16/0.16	0.35/0.32/0.25	0.60/0.41/0.37	0.10/0.09/0.07
50 × 5	3.85/3.57/3.43	2.27/1.65/2.18	1.79/1.39/1.13	2.52/2.05/1.98	2.99/2.61/2.42	1.41/1.13/1.04
50 × 10	3.24/3.15/3.01	1.81/1.96/1.68	1.49/1.24/1.08	2.15/1.81/1.62	2.44/2.23/2.12	1.33/1.17/0.92
50 × 20	2.57/2.49/2.43	1.93/1.61/1.69	1.33/1.07/0.89	1.65/1.42/1.28	2.34/2.06/2.03	1.16/0.93/0.82
100 × 5	4.64/4.06/3.98	2.64/2.35/2.34	2.23/1.72/1.38	4.37/4.11/3.95	2.93/2.67/2.33	1.51/1.27/1.09
100 × 10	3.61/3.24/3.07	2.20/1.82/1.52	1.84/1.53/1.21	3.44/3.19/3.10	2.69/2.23/2.13	1.37/1.04/0.88
100 × 20	2.96/2.71/2.51	2.00/1.66/1.54	1.73/1.35/1.03	2.87/2.66/2.45	2.38/2.01/1.82	1.29/0.96/0.81
200 × 10	3.95/3.64/3.49	1.98/1.71/1.35	1.88/1.43/1.21	3.62/3.48/3.37	2.59/2.19/1.90	1.33/0.88/0.63
200 × 20	3.04/2.82/2.67	1.62/1.34/1.19	1.61/1.17/1.02	2.89/2.78/2.64	2.07/1.77/1.51	1.10/0.74/0.53
500 × 20	2.14/2.09/2.07	1.29/0.99/0.76	1.23/0.96/0.79	2.00/2.00/2.00	1.79/1.47/1.28	0.86/0.50/0.31
Average	2.76/2.57/2.46	1.66/1.45/1.44	1.34/1.06/0.88	2.24/2.06/1.97	2.03/1.75/1.59	1.00/0.76/0.62

Each entry $x/y/z$ corresponds to the average percentage deviations for the setting of t to 30/60/90, respectively.

For the SDST-FSP- C_{\max} , the local search seems to play an important role. In fact, the ranking of the algorithms (leaving aside PACO for the moment) is roughly from those that give strong importance to the local search (IG_RS_{LS} and MA_{LS}), over less importance of the local search (MA) to those that do not use local search (IG_RS and GA). PACO, which uses the same local search as (IG_RS_{LS} and MA_{LS}) has a peculiar profile. While for SDST10 it yields on average better solutions than GA, IG_RS and MA, its deteriorates relatively to the others as the ratio of setup times to processing times increases, being even worse than IG_RS on the other instance sets. This behavior may be due to the type of information used when constructing solutions in PACO. Essentially, PACO exploits the information at which position a job is in a good sequence, while it does not consider the successor/predecessor relationship between jobs. However, as the relative influence of the setup times increases, this latter information becomes probably more important, which explains also the relatively poor performance of PACO.

While a considerable amount of testing has been carried out, comparing algorithms on the basis of means is rather weak. Hence, we have carried out an analysis of experiments using the ANOVA technique. In fact, we did four analyses, one for each instance set (SDST10, ..., SDST125), where the controlled factors are the type of instance, the algorithm applied and t . For each such experiment we consider all algorithms, values of t , instance types (we have 12 types, one for each combination of n and m , i.e., 1: 20 × 5, ..., 12: 500 × 20) and repetitions. The response variable is the relative percentage deviation or *RPD*. For each experiment all three hypotheses of the ANOVA technique (normality, homocedasticity and independence of the residuals) were checked and accepted. Of particular interest is the interaction between the algorithm and t factors. The means plot along with the Tukey confidence intervals (at the 95% confidence level) are given in Fig. 4 for experiment SDST10.

As we can see, for all algorithms, increasing the value of t results in better performance although

Table 2

Average percentage deviation from the best known solutions for the algorithms for the SDST-FSP- C_{\max} on instance sets SDST100 and SDST125 with the termination criteria set at $(n \times m/2) \times t$ milliseconds maximum CPU time

Instance	GA	MA	MA _{LS}	PACO	IG_RS	IG_RS _{LS}
<i>SDST100 instances</i>						
20 × 5	2.01/1.88/1.82	1.29/1.73/1.43	0.43/0.37/0.39	0.82/0.71/0.61	1.53/1.48/1.24	0.30/0.25/0.17
20 × 10	1.48/1.26/1.27	0.87/0.88/1.09	0.31/0.28/0.29	0.66/0.47/0.48	1.42/1.01/1.03	0.35/0.25/0.18
20 × 20	1.08/1.00/0.94	0.62/0.28/1.14	0.29/0.26/0.17	0.52/0.41/0.48	0.95/0.92/0.74	0.27/0.18/0.17
50 × 5	5.00/5.35/5.26	3.12/2.65/3.02	2.37/2.24/1.99	3.79/3.40/3.31	3.83/3.89/3.70	1.95/1.95/1.82
50 × 10	4.19/4.21/4.18	2.59/2.72/2.55	1.98/1.66/1.50	3.05/2.73/2.49	3.10/3.09/2.99	1.57/1.48/1.30
50 × 20	3.39/3.23/3.11	1.78/2.11/1.77	1.66/1.35/1.18	2.51/2.14/1.98	2.76/2.58/2.40	1.41/1.28/1.11
100 × 5	6.49/5.99/6.00	3.63/3.50/3.04	3.20/2.69/2.16	6.86/6.89/6.65	3.93/3.82/3.48	2.16/1.95/1.63
100 × 10	4.58/4.39/4.15	3.03/2.67/2.45	2.26/2.01/1.61	5.14/4.96/4.89	3.28/2.95/2.77	1.61/1.44/1.02
100 × 20	3.73/3.67/3.49	2.37/2.31/2.39	2.12/2.03/1.53	4.04/4.04/3.91	2.76/2.65/2.46	1.41/1.35/1.05
200 × 10	5.12/4.95/4.71	2.56/2.31/2.19	2.53/2.19/1.77	5.48/5.62/5.53	2.98/2.86/2.49	1.67/1.25/0.92
200 × 20	3.59/3.65/3.48	1.99/1.81/1.68	1.93/1.68/1.40	3.70/3.87/3.82	2.27/2.08/1.92	1.26/0.93/0.76
500 × 20	2.50/2.66/2.64	1.53/1.44/1.16	1.53/1.35/1.14	2.50/2.75/2.75	1.87/1.70/1.50	0.96/0.73/0.46
Average	3.60/3.52/3.42	2.11/2.03/1.99	1.72/1.51/1.26	3.26/3.17/3.07	2.56/2.42/2.23	1.24/1.09/0.88
<i>SDST125 instances</i>						
20 × 5	2.06/1.80/1.90	1.69/2.05/1.40	0.67/0.34/0.32	0.88/0.64/0.65	1.96/1.40/1.24	0.46/0.35/0.30
20 × 10	1.74/1.66/1.52	1.02/1.48/1.24	0.51/0.42/0.37	0.85/0.68/0.56	1.62/1.39/1.44	0.53/0.41/0.36
20 × 20	1.06/0.97/0.95	1.37/0.96/1.21	0.28/0.22/0.24	0.47/0.39/0.39	0.94/0.84/0.81	0.26/0.22/0.19
50 × 5	6.09/5.83/5.63	3.71/3.97/3.48	2.97/2.47/1.97	4.59/4.07/3.67	4.57/4.25/4.00	2.37/2.18/2.01
50 × 10	4.64/4.73/4.59	3.14/2.13/3.35	2.07/1.78/1.50	3.60/3.16/2.96	3.95/3.60/3.47	1.94/1.67/1.54
50 × 20	3.32/3.41/3.25	2.16/2.50/1.63	1.59/1.43/1.26	2.55/2.43/2.06	2.77/2.71/2.59	1.42/1.45/1.18
100 × 5	7.33/6.86/6.82	4.38/4.45/3.65	3.55/3.02/2.52	8.19/7.89/7.75	4.70/4.58/4.14	2.41/2.27/1.91
100 × 10	5.33/5.14/4.80	3.24/3.10/2.84	2.78/2.37/1.94	6.02/5.89/5.61	3.66/3.43/3.26	2.07/1.65/1.34
100 × 20	3.99/3.79/3.50	2.56/2.40/2.16	2.31/1.80/1.50	4.37/4.32/4.15	2.91/2.69/2.60	1.52/1.22/1.00
200 × 10	5.53/5.65/5.37	2.81/2.76/2.63	2.73/2.51/2.14	5.80/6.27/6.20	3.33/3.17/2.94	1.79/1.60/1.17
200 × 20	3.86/3.88/3.69	2.08/1.94/1.69	2.04/1.74/1.49	3.93/4.20/4.16	2.51/2.40/2.24	1.38/1.06/0.76
500 × 20	2.71/2.89/2.83	1.71/1.66/1.36	1.70/1.53/1.23	2.77/3.03/3.02	2.13/1.91/1.64	1.08/0.83/0.52
Average	3.97/3.88/3.74	2.49/2.45/2.22	1.93/1.64/1.37	3.67/3.58/3.43	2.92/2.70/2.53	1.44/1.24/1.02

Each entry $x/y/z$ corresponds to the average percentage deviations for the setting of t to 30/60/90, respectively.




Fig. 4. Plot of the average percentage deviation from best known solutions for the interaction between the type of algorithm and the different values of t for the stopping criterion. Instance set SDST10 and makespan criterion.

for $t = 60$ and $t = 90$ the differences are very small and even not statistically significant for MA. We can confirm also the previous observation that GA

is statistically worse than all other algorithms (recall that overlapping Tukey intervals means no statistically significant differences at the 95% confidence). We also observe that MA_{LS} is statistically better than MA, which is even true for MA_{LS} with $t = 30$ when compared to MA with $t = 90$. From the plot it is also clear that IG_RS is better than GA and also in this case we have that IG_RS with $t = 30$ yields better results than GA with $t = 90$. PACO in this experiment is comparable or marginally better than MA but the improved version MA_{LS} yields better results than PACO. Lastly, from the plot it is clear that IG_RS_{LS} is by far the highest performing algorithm and it beats MA_{LS} by a considerable margin. We can even see that IG_RS_{LS} with $t = 60$ gives statistically better results than MA_{LS} with $t = 90$.

In the previous analysis, we observed the results averaged across all instance types, but it is also possible to do a more in-depth analysis. In Fig. 5 we see for instance set SDST10 the interaction between the




Fig. 5. Plot of the average percentage deviation from best known solutions for the interaction between the type of algorithm and the type of instance. Instance set SDST10, $t = 90$ and makespan criterion.

type of instance and the algorithm factors after having fixed t to 90.

In this plot, we see that the range of the average percentage deviations is larger for some instances than on others. The instances with the largest deviations from the best known solutions are types 5, 6, 9 and 11, which correspond to those of size 50×10 , 50×20 , 100×20 and 200×20 . (The instances of these sizes are considered to be the hardest ones of the FSP- C_{\max} , from which they were derived.) It is especially on these instances that the differences among the algorithms are most pronounced and IG_RS_{LS} is clearly the best performing algorithm.

More or less the same conclusions as for instance set SDST10 can be drawn for the other three experiments with some minor exceptions. In Fig. 6 we




Fig. 6. Plot of the average percentage deviation from best known solutions for the interaction between the type of algorithm and the different values of t for the stopping criterion. Instance set SDST125 and makespan criterion.

show the interaction between the factors algorithm and t for SDST125.

As it has been mentioned, the differences among the different values of t are even more significant on this set (note that although in the graph the three lines appear to be closer together, the y-axis shows that the differences in the $AVRPD$ values are greater than in the SDST10 experiment). A striking difference from the SDST10 experiment is the strong deterioration of PACO (as also explained previously), which in this case is far worse than all algorithms except GA. This experiment also indicates that a more in-depth adaptation of PACO to the SDST-FSP- C_{\max} would be required.

As a first important observation we can say that the Iterated Greedy algorithms, especially the version with local search, are very robust since they resulted to be the best algorithms for the SDST-FSP- C_{\max} in all four instance sets and they were also state-of-the-art algorithms for the regular flowshop problem (see Ruiz and Stützle, 2007). This is very remarkable since the overall structure and parameter setting of the IG algorithm for the SDST flowshop remained essentially the same as for the FSP- C_{\max} and the only adaptations that were necessary were the implementation of the constructive heuristic, which was replaced by a known one for the SDST-FSP- C_{\max} , and of the local search algorithm.

4.2. Experimental results for the SDST-FSP-WT

Now we proceed to test the proposed algorithms as well as other existing methods for the total weighted tardiness objective. As a benchmark set we have chosen the already described SDST10, ..., SDST125 sets and augmented them by including weights and due dates. The four resulting benchmark sets are called DD_SDST10, ..., DD_SDST125. The weights for all n jobs (w_j) are drawn from a uniform $U[1, 10]$ distribution, which is also done for other benchmark sets for scheduling problems involving tardiness objectives (Congram et al., 2002). For obtaining the due dates we use an approach similar to that of Hasija and Rajendran (2004). In that work, the authors consider the regular flowshop with total tardiness objective and they assign tight due dates in their experiments. In order to determine the due dates with sequence dependent setup times, we apply the following steps:

1. Calculate the total processing time on all the m machines for each job:

$$\forall j, j \in N, \quad P_j = \sum_{i=1}^m p_{ij}$$

2. Calculate the average setup time for all possible following jobs and sum it for all machines:

$$\forall j, j \in N, \quad S_j = \sum_{i=1}^m \frac{\sum_{k=1, k \neq j}^n S_{ijk}}{n-1}$$

3. Assign a due date to each job:

$$\forall j, j \in N, \quad d_j = (P_j + S_j) \times (1 + \text{random} \times 3)$$

where random is a random number uniformly distributed in $[0, 1]$.

This method of generating due dates results in very tight to relatively tight due dates depending on the actual value of random for each job, i.e., if random is close to 0, then the due date of the job is going to be really tight as it would be more or less the sum of its processing times and average setup times on each machine. As a result, the job will have to be sequenced very early to avoid any tardiness. These 480 augmented instances as well as the best known solutions can be downloaded from <http://www.upv.es/gio/r Ruiz>.

As said in Section 2, we are only aware of three existing algorithms for the SDST-FSP-WT. We test the simulated annealing of Parthasarathy and Rajendran (1997a,b) which we refer to as SA_PR after the author's last names. We also consider the heuristic as well as the version with the improvement scheme proposed by Rajendran and Ziegler (1997, 2003). Following the nomenclature used by the authors, we refer to these two latter algorithms as HA and HA_IS2. We also test the proposed IG_RS and IG_RS_LS algorithms as well as the previously tested genetic algorithms GA, MA and the new MA_LS. We have refrained from including other methods such as PACO since it was proposed for a different objective and may require major adaptations.

It has to be mentioned that, for IG_RS, IG_RS_LS, GA, MA and MA_LS the initialization procedure is changed from NEH_RMB to NEH_EWDD. In order to validate this decision, we carry out an experiment in which the original NEH and the NEH_EWDD are compared. We also test the heuristic method HA. The results for all instance sets can be seen in Table 3.

We can see that the *AVRPD* values are very large for all heuristics, instance sets and sizes. From the

results, the heuristic HA, specifically designed for this problem and this objective, yields across all instances the worst results in most of the cases. When compared to the original NEH, HA obtains slightly better results only in instance sets DD_SDST100 and DD_SDST125. However, this poor performance of HA compared to NEH is mainly due to its very poor performance on the instances of size 20×20 , while on the other instance sizes it is occasionally slightly better than NEH. In any case, the modified NEH_EWDD gives much better results than either of HA or NEH: over the 480 instances, NEH_EWDD gives better result than HA in 420 cases and better results than the original NEH in 470 cases. Hence, there is no need of a statistical analysis to ascertain that NEH_EWDD is the best heuristic among these three. However, this performance comes at a computational cost. The CPU times measured on an Athlon XP 1600+ with 512 MBytes of main memory, given in Table 4, needed by NEH and NEH_EWDD are very large compared to those of HA. (Since the computation times do not depend on the relative duration of the setups to processing times and they are similar for the four groups of instances, Table 4 gives only the times for set DD_SDST125.)

As shown, HA is a very fast method, while both, NEH and NEH_EWDD, have very large CPU times especially for the largest instances. Given these CPU times, we might expect a degradation of the performance for the algorithms on the 500×20 instances. For the remaining cases, especially for cases where $n \leq 100$, the clear advantage of NEH_EWDD over HA would most probably prove beneficial since the CPU times needed in these cases are very small.

For testing the algorithms we use the same computer and methodology for the experiments as in the tests carried out for the SDST-FSP- C_{\max} . In this case, every algorithm is run 10 independent times on each instance for a maximum CPU time given by $(n \times m/2) \times 180$. Here we allow longer CPU times as for the SDST-FSP- C_{\max} , since calculating the total weighted tardiness objective is much more expensive than the C_{\max} objective and the initialization procedure requires large CPU times on the largest instances.

The results for instance sets DD_SDST10 and DD_SDST50 are given in Table 5, those for the sets DD_SDST100 and DD_SDST125 in Table 6.

When again first comparing the cross averages (results on the statistical significance are given

Table 3

Average relative percentage deviation from the best known solutions for the heuristic algorithms on the SDST-FSP-WT for all four instance sets

Instance	HA	NEH	NEH_EWDD	HA	NEH	NEH_EWDD
DD_SDST10 Instances			DD_SDST50 Instances			
20 × 5	78.72	83.00	25.53	55.31	47.74	22.51
20 × 10	676.49	650.10	218.41	454.86	486.43	169.72
20 × 20	1532.80	736.68	492.98	1571.25	1048.81	483.75
50 × 5	20.36	26.00	8.99	19.37	23.96	13.29
50 × 10	52.04	51.97	22.02	39.15	34.66	17.48
50 × 20	218.85	222.13	93.04	78.37	78.90	32.26
100 × 5	12.46	14.86	6.82	11.95	14.67	8.93
100 × 10	20.93	23.04	9.89	18.12	17.89	10.56
100 × 20	47.46	47.72	19.78	28.12	30.25	13.92
200 × 10	11.15	11.11	4.83	8.36	9.84	4.91
200 × 20	18.10	16.59	6.48	11.03	11.50	5.37
500 × 20	9.33	5.87	1.37	5.77	5.14	1.93
Average	224.89	157.42	75.84	191.80	150.82	65.39
DD_SDST100 Instances			DD_SDST125 Instances			
20 × 5	52.77	63.97	20.34	48.52	56.46	20.08
20 × 10	166.49	184.10	60.89	99.97	88.01	40.09
20 × 20	1130.55	1518.88	396.33	1522.99	1503.37	741.14
50 × 5	20.87	25.51	15.95	19.49	28.10	19.20
50 × 10	23.43	29.30	17.29	22.13	29.95	18.54
50 × 20	54.65	73.44	26.83	47.98	49.78	24.43
100 × 5	11.74	16.48	11.04	10.95	17.88	10.21
100 × 10	12.56	16.14	10.63	11.08	17.67	10.43
100 × 20	19.82	24.77	12.25	18.58	20.72	11.81
200 × 10	5.38	9.04	5.07	5.12	11.23	6.52
200 × 20	7.35	10.17	4.89	6.41	9.15	4.18
500 × 20	2.96	5.22	2.40	2.62	6.25	3.86
Average	125.72	164.75	48.66	151.32	153.21	75.87

Table 4

Average CPU times (in seconds) for the heuristic algorithms on the SDST-FSP-WT for instance set DD_SDST125

Instance	HA	NEH	NEH_EWDD
20 × 5	0.00	0.00	0.00
20 × 10	0.00	0.00	0.00
20 × 20	0.00	0.00	0.00
50 × 5	0.00	0.01	0.01
50 × 10	0.00	0.02	0.02
50 × 20	0.00	0.04	0.04
100 × 5	0.00	0.08	0.08
100 × 10	0.01	0.16	0.17
100 × 20	0.01	0.32	0.34
200 × 10	0.03	1.59	1.60
200 × 20	0.04	6.52	6.56
500 × 20	0.34	172.39	172.46
Average	0.04	15.09	15.11

below), we can see that IG_RS_{LS} is the best performing algorithm, followed by MA_{LS} and by IG_RS. The two worst performers are either SA_PR or HA_IS2, two algorithms that were specifically proposed for the SDST-FSP-WT. However,

it is worth mentioning that HA_IS2 gives the best results in all instance sets for the largest instances of 500 × 20. The previous analysis of the heuristics indicates that HA_IS2 benefits in this case from the fast HA initialization. (Additional experiments suggested that also the IG algorithms could benefit strongly on these largest instances from an initialization by HA.)

The performance of MA is roughly in the range of that of GA. Hence, the local search used in MA seems not to be very effective for the SDST-FSP-WT. Interestingly, the performance of MA is in some cases considerably worse than that of the GA method, as, for example, for the instances 20 × 20 in the group DD_SDST50. In fact, a more intensive local search than used in MA seems to be necessary, which is indicated by the much better performance of MA_{LS} when compared to MA. Given this importance of the local search, the high performance of IG_RS, which does not use any form of a descent local search, is remarkable: it has on all instances sets the third lowest average

Table 5

Average relative percentage deviation from the best known solutions for the algorithms for the SDST-FSP-WT on instance sets DD_SDST10 and DD_SDST50 with the termination criterion set at $(n \times m/2) \times 180$ milliseconds maximum CPU time

Instance	GA	MA	MA _{LS}	IG_RS	IG_RS _{LS}	SA_PR	HA_IS2
<i>DD_SDST10 instances</i>							
20 × 5	0.92	2.16	0.00	0.00	0.00	3.53	4.61
20 × 10	1.60	8.34	0.00	0.00	0.00	9.90	6.43
20 × 20	1.88	0.00	0.00	0.00	0.00	22.45	8.09
50 × 5	2.83	1.44	1.17	2.70	0.77	5.17	4.56
50 × 10	5.39	1.99	1.63	3.62	1.22	8.13	7.88
50 × 20	10.80	4.28	2.60	8.54	1.51	16.37	20.79
100 × 5	2.83	2.40	1.95	3.52	1.39	4.37	3.04
100 × 10	3.91	2.80	2.50	4.17	2.12	5.53	4.93
100 × 20	6.44	4.23	3.84	5.95	3.41	7.72	8.38
200 × 10	1.94	2.61	1.86	2.29	1.35	3.50	1.35
200 × 20	2.06	4.90	3.47	2.29	2.30	3.44	2.34
500 × 20	0.78	1.37	0.80	0.58	0.66	2.76	0.11
Average	3.45	3.04	1.65	2.80	1.23	7.74	6.04
<i>DD_SDST50 instances</i>							
20 × 5	1.45	6.96	0.00	0.02	0.00	5.92	7.20
20 × 10	5.28	13.87	0.09	0.15	0.15	12.58	24.90
20 × 20	0.43	30.78	0.00	0.00	0.00	4.79	1.68
50 × 5	5.42	2.19	2.00	3.77	1.79	7.59	7.31
50 × 10	5.83	2.02	2.24	3.79	2.41	8.42	9.13
50 × 20	6.75	2.53	2.08	4.69	2.80	10.95	12.97
100 × 5	4.47	3.14	2.66	4.82	2.14	6.70	4.19
100 × 10	4.65	3.09	2.61	4.59	2.51	6.65	6.25
100 × 20	5.42	3.39	3.15	5.22	3.29	7.12	6.06
200 × 10	1.95	2.56	1.71	2.37	1.07	3.65	1.12
200 × 20	1.61	4.10	2.70	1.95	1.61	3.10	1.51
500 × 20	1.43	1.93	1.42	1.04	0.95	3.76	0.00
Average	3.72	6.38	1.72	2.70	1.56	6.77	6.86

percentage deviations and on some specific instance sizes, it reaches better average solution quality than MA_{LS}.

We again carried out four analyses of experiments with the ANOVA technique in the same way as described before. However, in this case we controlled only two factors, namely the type of algorithm and the type of instance, since the parameter t is not used to vary the computation times.

As for the SDST-FSP- C_{\max} , also for the SDST-FSP-WT the overall results of the four experiments are rather similar and, hence, we illustrate the main findings giving some example results. Fig. 7 shows the Tukey means plot for the factor type of algorithm in the experiment DD_SDST100. We see that there are statistically significant differences among all algorithms with the only exception of MA and IG_RS that are, at the 95% confidence level, statistically equivalent. A finer examination and a closer look in Table 6 shows that there are differences depending on the instance type. Fig. 8 depicts a more precise example where we examine the

differences among the algorithms for a same instance group, here type 11, i.e., instances of size 200×20 .

In this case, we have that the differences observed in the $AVRPD$ values are statistically significant with the only exception of IG_RS_{LS} and HA_IS2, which from Table 6 have $AVRPD$ values of 1.60 and 1.61 respectively. Thanks to the number of replicates and the 480 instances, the Tukey intervals that result from the experiments are fairly narrow, which indicates that most of the observed differences in Tables 5 and 6 are statistically significant.

Similarly to the SDST-FSP- C_{\max} case, we have just examined the results averaged across all instance types. Fig. 9 shows, for set SDST10, the interaction between the type of instance and algorithm factors.

As shown in the results, the algorithms HA_IS2 and SA_PR have a poor performance for the smaller instances with improving results as the size of the instances grows.

Table 6

Average relative percentage deviation from the best known solutions for the algorithms for the SDST-FSP-WT on instance sets DD_SDST100 and DD_SDST125 with the termination criterion set at $(n \times m/2) \times 180$ milliseconds maximum CPU time

Instance	GA	MA	MA _{LS}	IG_RS	IG_RS _{LS}	SA_PR	HA_IS2
<i>DD_SDST100 instances</i>							
20 × 5	3.99	2.89	0.07	0.07	0.40	8.24	9.77
20 × 10	2.67	6.53	0.09	0.11	0.10	13.21	18.28
20 × 20	2.80	2.49	0.00	0.00	0.00	18.17	40.94
50 × 5	7.57	3.89	3.38	5.83	3.46	11.46	10.75
50 × 10	7.80	3.23	3.00	5.34	2.93	10.77	9.76
50 × 20	8.08	2.78	2.57	5.06	2.36	11.34	13.15
100 × 5	5.22	3.71	3.24	5.91	2.76	8.40	5.44
100 × 10	4.76	3.52	2.80	5.45	2.57	7.58	5.17
100 × 20	5.35	3.39	3.13	5.27	3.29	7.72	6.43
200 × 10	2.60	3.55	2.34	2.60	1.27	4.78	1.63
200 × 20	1.96	4.42	3.20	2.33	1.60	3.77	1.61
500 × 20	2.38	2.89	2.37	2.26	0.63	6.38	0.14
Average	4.60	3.61	2.18	3.35	1.78	9.32	10.26
<i>DD_SDST125 instances</i>							
20 × 5	3.68	5.54	0.41	0.15	0.40	10.22	14.82
20 × 10	3.91	6.28	0.10	0.00	0.04	10.80	15.12
20 × 20	3.01	11.57	0.00	0.00	0.00	42.38	14.08
50 × 5	8.36	4.07	3.75	7.20	3.62	12.67	11.43
50 × 10	8.01	3.63	3.49	6.25	3.79	11.16	10.50
50 × 20	7.83	2.63	2.75	5.56	2.74	11.45	10.70
100 × 5	5.01	3.73	2.87	5.58	2.01	9.27	4.78
100 × 10	4.69	3.38	2.81	5.30	2.52	7.88	5.35
100 × 20	4.68	3.07	2.61	4.64	2.79	6.97	5.97
200 × 10	3.10	4.38	2.97	3.46	0.99	6.00	1.23
200 × 20	2.08	3.48	3.10	1.84	1.17	4.10	1.35
500 × 20	3.20	3.86	3.16	3.13	0.45	7.43	0.32
Average	4.80	4.63	2.34	3.59	1.71	11.69	7.97




Fig. 7. Plot of the average percentage deviation from best known solutions for the type of algorithm factor. Instance set DD_SDST100 and total weighted tardiness criterion.

In summary, also for the SDST-FSP-WT, especially the proposed hybrid version of the Iterated Greedy algorithm gives excellent results. Only for the largest instances, another algorithm that is specifically designed for the SDST-FSP-WT gives




Fig. 8. Plot of the average percentage deviation from best known solutions for the type of algorithm factor. Instance set DD_SDST100, instance type 11 (200 × 20) and total weighted tardiness criterion.

slightly better results; however, this effect may be mainly due to the relatively slow initialization of our current algorithm and IG_RS_{LS} is clearly the best algorithm on all the other instances.

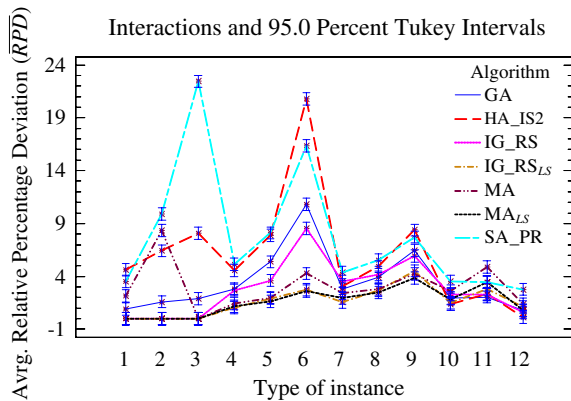


Fig. 9. Plot of the average percentage deviation from best known solutions for the interaction between the type of algorithm and the type of instance. Instance set DD_SDST10 and total weighted tardiness criterion.

5. Conclusion and future research

In this article, we have presented two new Iterated Greedy (IG) algorithms for SDST flowshop problems that were tackled under two of the most widely used and relevant objective functions, the minimization of the maximum completion time or makespan and the minimization of the total weighted tardiness. The first algorithm, called IG_RS, works by iterating over greedy construction heuristics, in this case based on the well-known NEH heuristic from Nawaz et al. (1983), and by applying two phases, named *destruction*, where some jobs are removed from the incumbent sequence, and *construction*, where the greedy heuristic is applied to reconstruct the sequence and to reinsert the jobs that were previously removed. The second, IG_RS_LS, is a straightforward extension of IG_RS by including an additional local search phase, in which a simple descent algorithm is applied. For both algorithms, extensive experiments and statistical analyses have been carried out.

The experimental results establish IG_RS_LS as a new state-of-the-art algorithm for the scheduling problems tackled, while IG_RS was shown to be the best algorithm that does not use an explicit local search phase. This excellent performance of the Iterated Greedy algorithms is very noteworthy because of two main reasons. Firstly, the Iterated Greedy algorithms have a very simple structure, they are easy to understand, and also very easy to implement. This makes them preferable over other, much more complex algorithms—even if the performance

would be the same. Additionally, the IG algorithms applied here are direct extensions of a successful IG algorithm for the FSP- C_{\max} Ruiz and Stützle (2007)—this is a first proof that IG algorithms may easily be extended to related problems maintaining their high performance. Second, these results provide some first evidence that the Iterated Greedy method is an excellent candidate for tackling complex scheduling environments as they are encountered in real-world environments.

There are mainly two interesting lines of future research, those concerning algorithmic issue and more realistic scheduling environments, respectively. In the first direction, it would be interesting to examine for which classes of combinatorial problems the IG method can be highly competitive. Certainly, the availability of high-performing construction methods are an essential ingredient, but also other factors may be central to the performance of IG algorithms. In the second direction, we would like to test IG based algorithms on other features of flowshop scheduling environments like no-wait, availability of parallel machines, other objectives, etc. This direction is important to identify with which characteristics of scheduling problems the IG algorithms can deal effectively. The ultimate goal is to apply IG to more complex scheduling environments, both with the consideration of more realistic constraints as well as with other objectives. Clear examples of this are flowshops or jobshops where the production floor is divided up into stages and at each stage several identical or even unrelated parallel machines are present. These environments, the so-called hybrid shops, are of great importance in practice and IG methods may be excellent candidates for tackling them.

Acknowledgements

Rubén Ruiz is partly funded by the Spanish Department of Science and Technology (research project ref. DPI2004-06366-C03-01). Thomas Stützle acknowledges support of the Belgian FNRS, of which he is a research associate.

References

- Allahverdi, A., Ng, C.T., Cheng, T.C.E., Kovalyov, M.Y., in press. A survey of scheduling problems with setup times or costs. *European Journal of Operational Research*, doi:10.1016/j.ejor.2006.06.060.
- Cesta, A., Oddi, A., Smith, S.F., 2000. Iterative flattening: A scalable method for solving multi-capacity scheduling prob-

- lems. In: Proceedings of the Seventeenth National Conference on Artificial Intelligence. AAAI Press/The MIT Press, Menlo Park, CA, USA, pp. 742–747.
- Congram, R.K., Potts, C.N., van de Velde, S., 2002. An iterated dynasearch algorithm for the single-machine total weighted tardiness scheduling problem. *INFORMS Journal on Computing* 14 (1), 52–67.
- Corwin, B.D., Esogbue, A.O., 1974. Two machine flow shop scheduling problems with sequence dependent setup times: A dynamic-programming approach. *Naval Research Logistics* 21 (3), 515–524.
- Das, S.R., Gupta, J.N.D., Khumawala, B.M., 1995. A savings index heuristic algorithm for flowshop scheduling with sequence-dependent set-up times. *Journal of the Operational Research Society* 46 (11), 1365–1373.
- Gupta, J.N.D., 1986. Flowshop schedules with sequence dependent setup times. *Journal of the Operations Research Society of Japan* 29 (3), 206–219.
- Gupta, J.N.D., Darrow, W.P., 1986. The two-machine sequence dependent flowshop scheduling problem. *European Journal of Operational Research* 24 (3), 439–446.
- Hasija, S., Rajendran, C., 2004. Scheduling in flowshops to minimize total tardiness of jobs. *International Journal of Production Research* 42 (11), 2289–2301.
- Hoos, H.H., Stützle, T., 2004. *Stochastic Local Search: Foundations and Applications*. Morgan Kaufmann Publishers, San Francisco, CA.
- Jacobs, L.W., Brusco, M.J., 1995. A local-search heuristic for large set-covering problems. *Naval Research Logistics* 42 (7), 1129–1140.
- Laurent, M., Van Hentenryck, P., 2004. Iterative relaxations for iterative flattening in cumulative scheduling. In: Zilberstein, J., Koehler, J., Koenig, S. (Eds.), *Proceedings of the Fourteenth International Conference on Automated Planning and Scheduling (ICAPS 2004)*. AAAI Press/The MIT Press, Menlo Park, CA, USA, pp. 200–208.
- Lourenço, H.R., Martin, O.C., Stützle, T., 2003. Iterated local search. In: Glover, F., Kochenberger, G.A. (Eds.), *Handbook of Metaheuristics*. Kluwer Academic Publishers, Boston, pp. 321–353.
- Marchiori, E., Steenbeek, A., 2000. An evolutionary algorithm for large scale set covering problems with application to airline crew scheduling. In: Cagnoni, S. et al. (Eds.), *Real-World Applications of Evolutionary Computing, EvoWorkshops 2000, Lecture Notes in Computer Science*, vol. 1803. Springer-Verlag, Berlin, Germany, pp. 367–381.
- Nawaz, M., Ensore Jr., E.E., Ham, I., 1983. A heuristic algorithm for the m machine, n job flowshop sequencing problem. *Omega—International Journal of Management Science* 11 (1), 91–95.
- Osman, I.H., Potts, C.N., 1989. Simulated annealing for permutation flowshop scheduling. *Omega—International Journal of Management Science* 17 (6), 551–557.
- Parthasarathy, S., Rajendran, C., 1997a. An experimental evaluation of heuristics for scheduling in a real-life flowshop with sequence-dependent setup times of jobs. *International Journal of Production Economics* 49 (3), 255–263.
- Parthasarathy, S., Rajendran, C., 1997b. A simulated annealing heuristic for scheduling to minimize mean weighted tardiness in a flowshop with sequence-dependent setup times of jobs – a case study. *Production Planning & Control* 8 (5), 475–483.
- Pinedo, M., 2002. *Scheduling: Theory, Algorithms, and Systems*, second ed. Prentice-Hall, Upper Saddle, NJ.
- Rajendran, C., Ziegler, H., 1997. A heuristic for scheduling to minimize the sum of weighted flowtime of jobs in a flowshop with sequence-dependent setup times of jobs. *Computers & Industrial Engineering* 33 (1–2), 281–284.
- Rajendran, C., Ziegler, H., 2003. Scheduling to minimize the sum of weighted flowtime and weighted tardiness of jobs in a flowshop with sequence-dependent setup times. *European Journal of Operational Research* 149 (3), 513–522.
- Rajendran, C., Ziegler, H., 2004. Ant-colony algorithms for permutation flowshop scheduling to minimize makespan/total flowtime of jobs. *European Journal of Operational Research* 155 (2), 426–438.
- Ríos-Mercado, R.Z., Bard, J.F., 1998a. Computational experience with a branch-and-cut algorithm for flowshop scheduling with setups. *Computers & Operations Research* 25 (5), 351–366.
- Ríos-Mercado, R.Z., Bard, J.F., 1998b. Heuristics for the flow line problem with setup costs. *European Journal of Operational Research* 110 (1), 76–98.
- Ríos-Mercado, R.Z., Bard, J.F., 1999a. A branch-and-bound algorithm for permutation flow shops with sequence-dependent setup times. *IIE Transactions* 31 (8), 721–731.
- Ríos-Mercado, R.Z., Bard, J.F., 1999b. An enhanced TSP-based heuristic for makespan minimization in a flow shop with setup times. *Journal of Heuristics* 5 (1), 53–70.
- Ríos-Mercado, R.Z., Bard, J.F., 2003. The flow shop scheduling polyhedron with setup times. *Journal of Combinatorial Optimization* 7 (3), 291–318.
- Ruiz, R., Maroto, C., 2005. A comprehensive review and evaluation of permutation flowshop heuristics. *European Journal of Operational Research* 165 (2), 479–494.
- Ruiz, R., Stützle, T., 2007. A simple and effective iterated greedy algorithm for the permutation flowshop scheduling problem. *European Journal of Operational Research* 177 (3), 2033–2049.
- Ruiz, R., Maroto, C., Alcaraz, J., 2005. Solving the flowshop scheduling problem with sequence dependent setup times using advanced metaheuristics. *European Journal of Operational Research* 165 (1), 34–54.
- Schrumpf, G., Schneider, J., Stamm-Wilbrandt, H., Dueck, G., 2000. Record breaking optimization results using the ruin and recreate principle. *Journal of Computational Physics* 159, 139–171.
- Simons Jr., J.V., 1992. Heuristics in flow shop scheduling with sequence dependent setup times. *Omega—International Journal of Management Science* 20 (2), 215–225.
- Srikar, B.N., Ghosh, S., 1986. A MILP model for the n -job, m -stage flowshop with sequence dependent set-up times. *International Journal of Production Research* 24 (6), 1459–1474.
- Stafford Jr., E.F., Tseng, F.T., 1990. On the Srikar-Ghosh MILP model for the $n \times m$ SDST flowshop problem. *International Journal of Production Research* 28 (10), 1817–1830.
- Stafford Jr., E.F., Tseng, F.T., 2002. Two models for a family of flowshop sequencing problems. *European Journal of Operational Research* 142 (2), 282–293.

- Stützle, T., 1998. Applying iterated local search to the permutation flow shop problem. Technical Report AIDA-98-04, FG Intellektik, FB Informatik, TU Darmstadt.
- Taillard, E., 1990. Some efficient heuristic methods for the flowshop sequencing problem. *European Journal of Operational Research* 47 (1), 65–74.
- Taillard, E., 1993. Benchmarks for basic scheduling problems. *European Journal of Operational Research* 64 (2), 278–285.
- Tseng, F.T., Stafford Jr., E.F., 2001. Two MILP models for the $n \times m$ SDST flowshop sequencing problem. *International Journal of Production Research* 39 (8), 1777–1809.